




Netflix Official Final Report

Members: Sha Sha, Yu Wu, Diyi Liu,
Haotian Sheng, Wen Gu



Overview

- Final Performance
- Restricted Boltzmann Machine
- Singular Value Decomposition
- K-Nearest Neighbor
- Blending
- Division of Labor

Final Performance

- Test Performance:

- RMSE: 0.88136
- Above water: 7.3618%

- Quiz Performance:

- RMSE: 0.88051
- Above water: 7.4511% (Indicating overfitting on quiz)

- Novelty Score:

- 0.4294

Models

Basic Restricted Boltzmann Machine

- Two RBM models:
 - Model 1 trained with base, validation, hidden
 - Model 2 trained with base, validation, hidden, probe
- Quiz Performance:
 - Model 1: 3.02%
 - Model 2: 3.52%
 - Indicates **shift between training set and test set**
- Model Detail:
 - 88,850 visible units (known), 100 hidden units(need to be tuned)
 - Total number of parameters: 8,973,950

Basic Restricted Boltzmann Machine -- Hyperparameter Tuning

- Number of hidden units:
 - We also tried 150, 200 hidden units, the performance improved is very tiny.
 - Linear relationship between number of hidden units and training time.
- Contrastive Divergence (CD):
 - Adaptive Gibbs Sampling Steps (T)
 - Epoch 0: $T = 1$, Epoch 19: $T = 3$, Epoch 28: $T = 5$, Epoch 43: $T = 9$
 - Linear relationship between steps and training time.
- Adaptive Learning Rate:
 - Start with $2e-5/\text{user}$, weight decay = 0.999, slow convergence
 - At epoch 46, increase the learning rate to $3e-5/\text{user}$, weight decay = 0.999
- Epochs: 150

Epoch 1

Random Initialized RBM Model

Learning Rate $2e-5$ /user, weight decay = 0.999

Training Setup:

AWS c5n.2xlarge, 3.0
Ghz Intel Xeon Platinum
8000 series CPU, 21GiB
RAM (single thread)

Training Dataset:
base data



Epoch 46


Pretrained RBM Model

Learning Rate $3e-5$ /user

weight decay = 0.999

Dataset:
base,
validation,
hidden

Dataset:
base,
validation,
hidden,
probe



Epoch 150

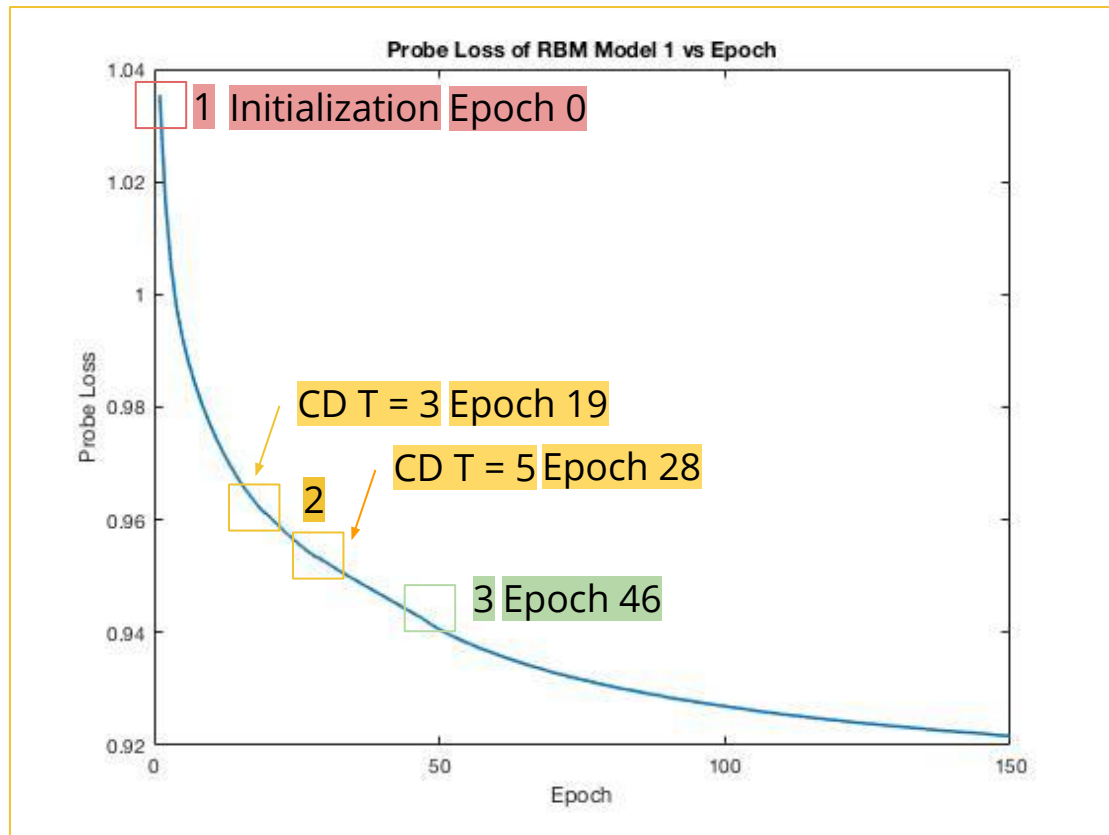
RBM Model 1

Used for blending training

RBM Model 2

For final blending

Basic Restricted Boltzmann Machine -- Probe Loss Function



1: Initialization does matter

- Weight / Hidden Units bias \sim Normal(0, 0.01)
- Visible Units bias $\log(p/(1-p))$
- First epoch RMSE loss decrease by 0.03

2: Increasing CD steps in training helps faster convergence

3: Feed in more data (Base, Hidden, Validation); Adjust Learning rate.

Basic Restricted Boltzmann Machine -- Optimization

- Computational Challenge
 - Large number of parameters embedded in the model
- Sparse Matrix
 - Missing ratings do not need to be considered in the sampling and update
 - Employ Unordered_map to lock on to the weights of presenting ratings
 - By omitting the missing weights, we reduce the computation by 5x
- Eigen Linear Solver
 - The Eigen Linear solver is 1.5x slower
- Other tricks
 - Using g++ compiler optimization (g++ -O2) reduce running time by 7x
 - -O2 is faster than -O1, but almost the same as -O3
- Training Time/ Epoch : ~43mins, CD step = 1

Singular Value Decomposition

- Best Performance of Single Model:

- More advanced SVD: (~2mins/epoch)

- $\operatorname{argmin}_{U,V,a,b} \frac{\lambda}{2} (\|U\|^2 + \|V\|^2 + \|a\|^2 + \|b\|^2) + \sum_{(i,j) \in \mathcal{S}} \left((Y_{i,j} - \mu) - (u_i^T v_j + a_i + b_j) \right)^2$

- 4.33% (K = 20)


- SVD ++ (~15mins/epoch)

- 5.36% (K = 80)

- Time SVD ++ (~20mins/epoch)

- 6.0563% (K= 80)

Singular Value Decomposition



Singular Value Decomposition

- Implementation & Optimization

- parameters referred to Koren's paper: "Advances in Collaborative Filtering"
- learning rate decay per iteration (*0.9)
- update y per user to speed up (tradeoff between accuracy and speed)

- SVD++

for *iter* \leftarrow 0 to *maxIter*

for *data_point* in *training_dataset*

if *new_user* then

curr_user \leftarrow *new_user*

update *sumYj* for *curr_user*

perform updates for *Bu, Bi, Pu, Qi*

(save *tempYPart* for the update of y)

if (*last_data_point* of *curr_user*) then

update *y* with *tempYPart*

update *sumYj* for *curr_user*

per data point

per user

(approximation to speed up)

$$\begin{aligned}
 & \text{0.007} \quad \text{0.005} \\
 b_u & \leftarrow b_u + \boxed{\gamma} (e_{ui} - \boxed{\lambda_5} \cdot b_u) \\
 b_i & \leftarrow b_i + \gamma \cdot (e_{ui} - \lambda_5 \cdot b_i) \\
 q_i & \leftarrow q_i + \gamma \cdot (e_{ui} \cdot (p_u + |\mathcal{N}(u)|^{-\frac{1}{2}} \boxed{\text{sumYj}}) - \lambda_6 \cdot q_i) \\
 p_u & \leftarrow p_u + \gamma \cdot (e_{ui} \cdot q_i - \lambda_6 \cdot p_u) \\
 \forall j \in \mathcal{R}(u) : & \quad \text{tempYPart} \quad \text{0.015} \\
 y_j & \leftarrow y_j + \gamma \cdot (\boxed{e_{ui} \cdot |\mathcal{R}(u)|^{-\frac{1}{2}} \cdot q_i} - \boxed{\lambda_6} \cdot y_j)
 \end{aligned}$$

Singular Value Decomposition

- Time SVD++

for *iter* \leftarrow 0 to *maxIter*
 for *data_point* in *training_dataset*
 if *new_user* then
 curr_user \leftarrow *new_user*
 update *sumYj* for *curr_user*
 perform updates for *B_u*, *B_i*, *B_{u,t}*, *P_u*, *Q_i*, *B_{i, Bin}*, *α_u*
 (save *tempYPart* for the update of *y*)
 if (last data point of *curr_user*) then
 update *y* with *tempYPart*
 update *sumYj* for *curr_user*

per data point \leftarrow perform updates for *B_u*, *B_i*, *B_{u,t}*, *P_u*, *Q_i*, *B_{i, Bin}*, *α_u*

per user (approximation to speed up) \rightarrow update *y* with *tempYPart*
 update *sumYj* for *curr_user*

Other Speed Up:

$$dev_u(t) = sign(t - t_u) \cdot |t - t_u|^\beta$$

- Calculate t_u for each user in initialization for $dev_u(t_{ui})$
- Calculate $dev_u(t_{ui})$ in advance, save in $Dev[user][date]$ matrix for future indexing

Differences from SVD++ in Time SVD++

$$b_u \leftarrow b_u + \gamma (e_{ui} - \lambda_5 \cdot b_u)$$

$$b_i \leftarrow b_i + \gamma \cdot (e_{ui} - \lambda_5 \cdot b_i)$$

$$b_{u,t} \leftarrow b_{u,t} + \gamma \cdot (e_{ui} - \lambda_5 \cdot b_{u,t})$$

$$q_i \leftarrow q_i + \gamma \cdot (e_{ui} \cdot (p_u + |R(u)|^{-\frac{1}{2}} \sum_{j \in R(u)} y_j) - \lambda_6 \cdot q_i)$$

$$p_u \leftarrow p_u + \gamma \cdot (e_{ui} \cdot q_i - \lambda_6 \cdot p_u)$$

$$b_{i, Bin(t_{ui})} \leftarrow b_{i, Bin(t_{ui})} + \gamma \cdot (e_{ui} - \lambda_5 \cdot b_{i, Bin(t_{ui})})$$

$$\alpha_u \leftarrow \alpha_u + \gamma_2 \cdot (e_{ui} \cdot dev_u(t_{ui}) - \lambda_5 \cdot \alpha_u)$$

$$\forall j \in R(u): tempYPart$$

$$y_j \leftarrow y_j + \gamma \cdot (e_{ui} \cdot |R(u)|^{-\frac{1}{2}} \cdot q_i - \lambda_6 \cdot y_j)$$

K-Nearest Neighbor

- Implementation

- Pearson Calculation:

Struct: *Pearson Intermediates* Array [# movies]

$$r_{xy} = \frac{n \sum x_i y_i - \sum x_i \sum y_i}{\sqrt{n \sum x_i^2 - (\sum x_i)^2} \sqrt{n \sum y_i^2 - (\sum y_i)^2}}$$


- Prediction: weighted average rating

$$r_{ui} = \frac{\sum_{j=1}^k p_{ij} \cdot (r_{uj} - \bar{r}_j)}{\sum_{j=1}^k |p_{ij}|} + \bar{r}_i$$

top k movies similar to movie i rated by user u
 movie j rating given by user u
 average rating of movie j given by all users
 average rating of movie i given by all users
 new test point: user u, movie i
 Pearson correlation coefficient between movie i and j

K-Nearest Neighbor

- Performance
 - not very well (below water) as a single model
 - $k = 8$: RMSE = 0.9844
 - $k = 16$: RMSE = 0.967442
 - $k = 24$: RMSE = 0.964737
 - $k = 32$: RMSE = 0.964719
 - $k = 40$: RMSE = 0.965459
 - $k = 48$: RMSE = 0.966065



Blending

Blending

- Data

- training data:
 - X_{train} = probe predictions (trained on training data)
 - y_{train} = probe true ratings
 - RBM
 - SVD (advanced SVD, SVD ++, Time SVD ++)
 - different number of factors ($K = 20, 50, 80, 200, 300$)
 - KNN
 - different number of nearest neighbors ($K = 8, 16, 32$)
- test data:
 - X_{test} = qual predictions (trained on training data + probe data)

Blending

- Models

- Referred to *“Combining Predictions for Accurate Recommender Systems”*
- Generally, more models, better performance
- Linear Regression:
 - Ridge Regression
 - regularization strength $\alpha=5e-6$
 - Best performance: ~7.36%
- XGBoost: (optimized distributed gradient boosting)
 - XGBoost Regressor
 - $max_depth=6$, $n_estimators=70$, other parameters set default
 - Best performance: ~7.4%
- Neural Network:
 - Two dense hidden layers with 64 units, “relu” activation
 - Optimizer: Stochastic gradient descent (SGD), $lr=5e-4$, $decay=5e-7$
 - Best performance: ~7.4%

Blending

- Further combination: Mean
 - Calculate the mean predictions of the blended qual predictions
 - Ridge Regression
 - XGBoost Regressor
 - Nerual Network
 - Performance:
 - input: all three prediction data
 - 7.42%
 - input: exclude predictions from Ridge Regression
 - 7.45%

Division of Labor

- RBM: Wen Gu, Haotian Sheng
- SVD: Yu Wu, Sha Sha, Diyi Liu
- KNN: Yu Wu, Sha Sha, Diyi Liu
- Blending: Yu Wu, Sha Sha, Diyi Liu



Questions?





Thanks!

